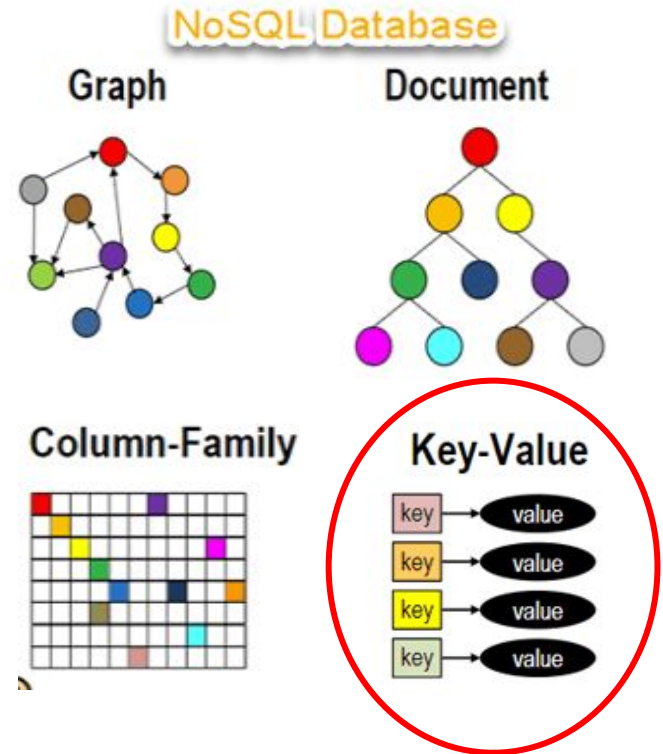


Assignment 6 NoSQL

Computing Lab - II (CS69012)
IIT Kharagpur

What is NoSQL?


- "Not Only SQL" or "Not SQL" is a non-relational Database Management System, that does not require a fixed schema, avoids joins, and is easy to scale
- Used for distributed data stores with humongous data storage needs (Big data and real-time web apps)
- Can store structured, semi-structured, unstructured and polymorphic data




Why NoSQL?

- System response time becomes slow when RDBMS is used for massive volumes of data
- To resolve this problem, we could "scale up" our systems by upgrading existing hardware, which is expensive
- The alternative for this problem is to distribute database load on multiple hosts whenever the load increases. This method is known as "scaling out"
- NoSQL database is non-relational, so it scales out better than relational databases

Scale-Up (*vertical scaling*):



Scale-Out (*horizontal scaling*):





Understanding Map-Reduce

Map:

- Grab the relevant data from the source.
- User function gets called for each chunk of input.
- Spit out (key, value) pairs.

Reduce:

- Aggregate the results.
- User function gets called for each unique key with all values corresponding to that key.

Map-Reduce: What happens in between?

- Map

- Grab the relevant data from the source (parse into key, value)
- Write it to an intermediate file

- Partition

- Partitioning: identify which of R reducers will handle which keys
- Map partitions data to target it to one of R Reduce workers based on a partitioning function (both R and partitioning function user defined)

Map Worker


- Shuffle & Sort

- Shuffle: Fetch the relevant partition of the output from all mappers
- Sort by keys (different mappers may have sent data with the same key)


- Reduce

- Input is the sorted output of mappers
- Call the user *Reduce* function per key with the list of values for that key to aggregate the results


Reduce Worker



Mapper Phase



Reducer Phase



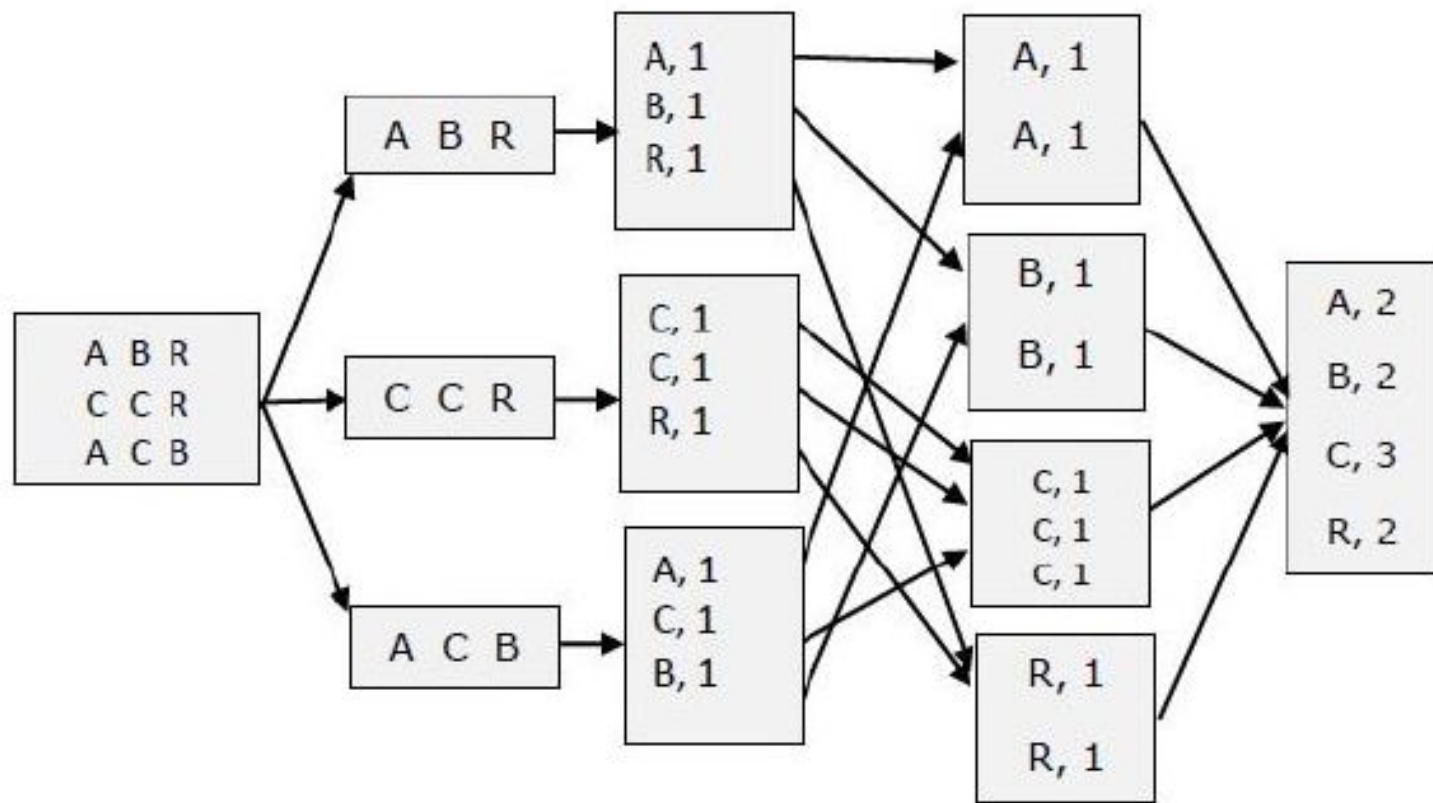
Input

Split

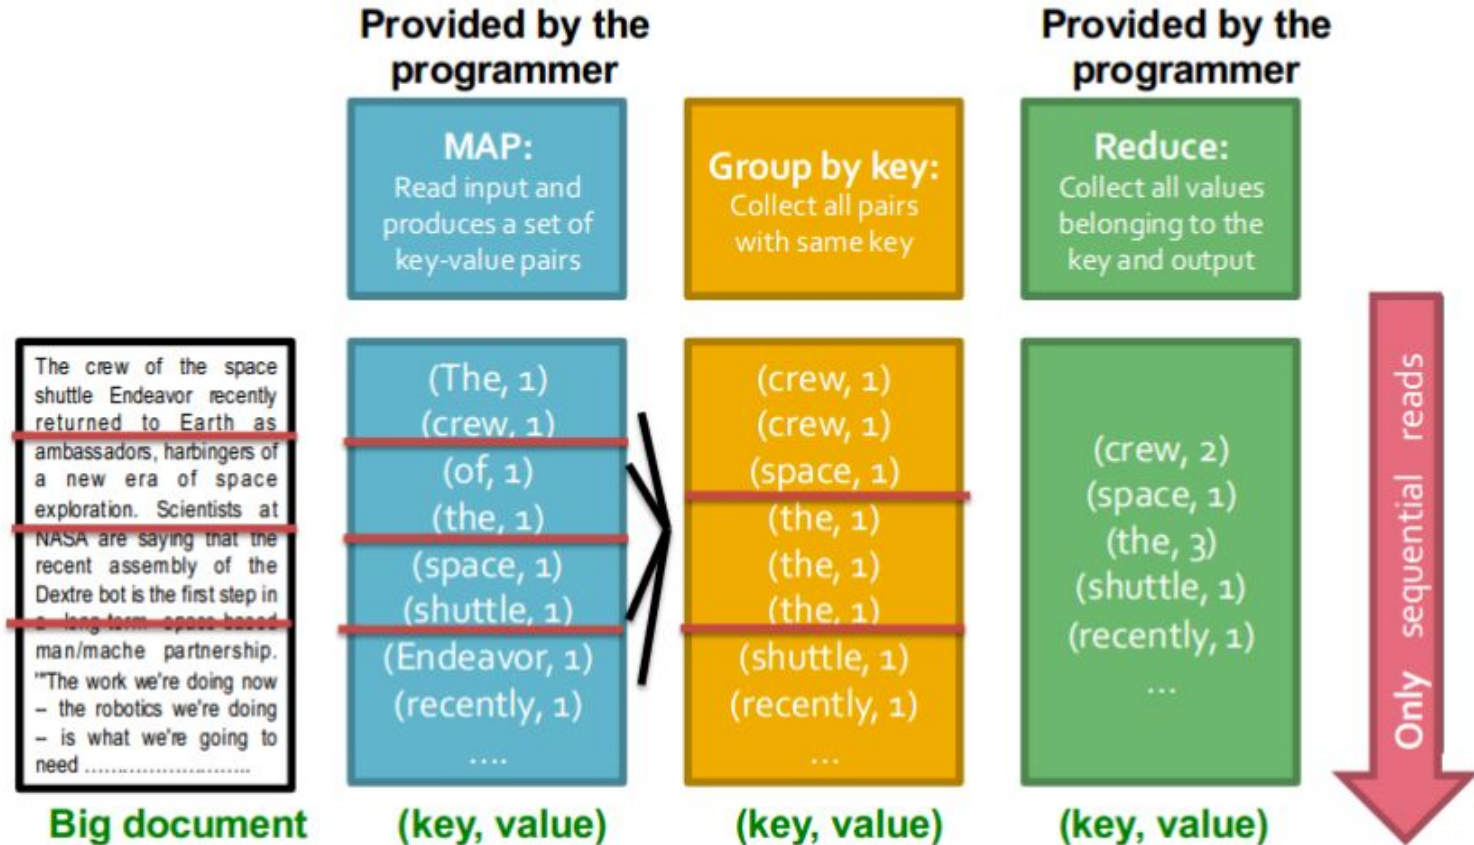
Map Phase

Shuffle and
Sort

Reduce
Phase



Map-Reduce Example



Sample Code (Word Count):

```
import sys
```

```
# input comes from STDIN (standard input)
```

```
for line in sys.stdin:
```

```
    # remove leading and trailing whitespace
```

```
    line = line.strip()
```

```
    # split the line into words
```

```
    words = line.split()
```

```
    # increase counters
```

```
    for word in words:
```

```
        # write the results to STDOUT (standard output);
```

```
        # what we output here will be the input for the
```

```
        # Reduce step, i.e. the input for reducer.py
```

```
        #
```

```
        # tab-delimited; the trivial word count is 1
```

```
        print '%s\t%s' % (word, 1)
```

Mapper.py

Sort

```
from operator import itemgetter
```

```
import sys
```

```
current_word = None
```

```
current_count = 0
```

```
word = None
```

```
# input comes from STDIN
```

```
for line in sys.stdin:
```

```
    # remove leading and trailing whitespace
```

```
    line = line.strip()
```

```
    # parse the input we got from mapper.py
```

```
    word, count = line.split('\t', 1)
```

```
    # convert count (currently a string) to int
```

```
    try:
```

```
        count = int(count)
```

```
    except ValueError:
```

```
        continue
```

```
# this IF-switch only works because Hadoop sorts map output
```

```
# by key (here: word) before it is passed to the reducer
```

```
if current_word == word:
```

```
    current_count += count
```

```
else:
```

```
    if current_word:
```

```
        # write result to STDOUT
```

```
        print '%s\t%s' % (current_word, current_count)
```

```
    current_count = count
```

```
    current_word = word
```

```
# do not forget to output the last word if needed!
```

```
if current_word == word:
```

```
    print '%s\t%s' % (current_word, current_count)
```

Reducer.py



Understanding the Assignment

Data Set:

- The dataset describes a **Facebook network**.
- There is only 1 file (listing all edges).
- Node: a facebook user
Edge: friendship
- Edges are undirected.

node_1,node_2

0,1838

0,1744

0,14

0,2543

1,1009

1,1171

1,1465

1,2069

1,2080

1,1856

1,3799

1,1033

1,2717

1,300

1,1603

1,942

1,3678

1,952

Example Query

- Find the users with more than 10 friends.
- Write a script mapper.py which takes data as input and transforms it into key-value pairs. (here <node-id,1>)
- Print the key-value pairs and use sort function to arrange it.
- The sorted key-value pairs are then passed to the reducer to group by key and perform necessary actions.
- Here reducer should count the degree for each particular node and if it is greater than 10, then it should print that node-id or the user.

***Since we do not have access to a hadoop cluster, we will be testing our codes on a linux system as follows:**

```
cat input_file | python mapper.py | sort | python reducer.py
```

```
python mapper.py | sort | python reducer.py
```

Deliverables

- For each of the query, make a directory with name “Query<no.>”
- Each directory should contain mapper + reducer + result + Makefile.
- In total, your submission should have
 - 4 mapper.py
 - 4 reducer.py
 - 4 result.txt
 - 4 Makefile
 - 1 readme



Thank you.

